Harvard PH125.9x Capstone: Mushroom Classification Analysis

Mauro Berlanda

April 2020

Introduction

The last assignment of the Datascience Professional Certificate by HarvardX on edx is submitting its own report. The main goal of the project is to prove the ability to clearly communicate the process and the insights gained from an analysis.

We are going to use for this analysis the Mushroom records drawn from The Audubon Society Field Guide to North American Mushrooms (1981. This data set includes descriptions of hypothetical samples corresponding to 23 species of gilled mushrooms in the Agaricus and Lepiota Family.

> The csv file containing the data was originally downloaded from Kaggle due to its ease of manipulation. The file has been commited in a github repository since Kaggle downloads require authentication. Being unable to retrieve the raw zip file due to a corrupted output (`unzip error -1`), my script is downloading the uncompressed csv file. It does not exceed 365Kb, so it can be requested without any performance or network traffic concern.

```r
file_url <- "https://raw.githubusercontent.com/mberlanda/ph125-9x-data-science-capstone/master/mushroom
csv_filepath <- "./mushrooms.csv"

# Download the csv file if needed
if (!file.exists(csv_filepath)) {
  download.file(file_url, csv_filepath)
}

# Use read.csv to parse the file converting strings to factors
mushrooms <- read.csv(csv_filepath, header=TRUE, sep=",", stringsAsFactors=TRUE)
# Explore the columns and types of the dataset
str(mushrooms)
```

```
## 'data.frame':    8124 obs. of  23 variables:
##  $ class                   : Factor w/ 2 levels "e","p": 2 1 1 2 1 1 1 1 2 1 ...
##  $ cap.shape               : Factor w/ 6 levels "b","c","f","k",..: 6 6 1 6 6 6 1 1 6 1 ...
##  $ cap.surface             : Factor w/ 4 levels "f","g","s","y": 3 3 3 4 3 4 3 4 4 3 ...
##  $ cap.color               : Factor w/ 10 levels "b","c","e","g",..: 5 10 9 9 4 10 9 9 9 10 ...
##  $ bruises                 : Factor w/ 2 levels "f","t": 2 2 2 2 1 2 2 2 2 2 ...
##  $ odor                    : Factor w/ 9 levels "a","c","f","l",..: 7 1 4 7 6 1 1 4 7 1 ...
##  $ gill.attachment         : Factor w/ 2 levels "a","f": 2 2 2 2 2 2 2 2 2 2 ...
##  $ gill.spacing            : Factor w/ 2 levels "c","w": 1 1 1 1 2 1 1 1 1 1 ...
##  $ gill.size               : Factor w/ 2 levels "b","n": 2 1 1 2 1 1 1 1 1 2 1 ...
##  $ gill.color              : Factor w/ 12 levels "b","e","g","h",..: 5 5 6 6 5 6 3 6 8 3 ...
##  $ stalk.shape             : Factor w/ 2 levels "e","t": 1 1 1 1 2 1 1 1 1 1 ...
##  $ stalk.root              : Factor w/ 5 levels "?","b","c","e",..: 4 3 3 4 4 3 3 3 4 3 ...
##  $ stalk.surface.above.ring: Factor w/ 4 levels "f","k","s","y": 3 3 3 3 3 3 3 3 3 3 ...
##  $ stalk.surface.below.ring: Factor w/ 4 levels "f","k","s","y": 3 3 3 3 3 3 3 3 3 3 ...
```

```
##  $ stalk.color.above.ring : Factor w/ 9 levels "b","c","e","g",..: 8 8 8 8 8 8 8 8 8 8 ...
##  $ stalk.color.below.ring : Factor w/ 9 levels "b","c","e","g",..: 8 8 8 8 8 8 8 8 8 8 ...
##  $ veil.type              : Factor w/ 1 level "p": 1 1 1 1 1 1 1 1 1 1 ...
##  $ veil.color             : Factor w/ 4 levels "n","o","w","y": 3 3 3 3 3 3 3 3 3 3 ...
##  $ ring.number            : Factor w/ 3 levels "n","o","t": 2 2 2 2 2 2 2 2 2 2 ...
##  $ ring.type              : Factor w/ 5 levels "e","f","l","n",..: 5 5 5 5 1 5 5 5 5 5 ...
##  $ spore.print.color      : Factor w/ 9 levels "b","h","k","n",..: 3 4 4 3 4 3 3 4 3 3 ...
##  $ population             : Factor w/ 6 levels "a","c","n","s",..: 4 3 3 4 1 3 3 4 5 4 ...
##  $ habitat                : Factor w/ 7 levels "d","g","l","m",..: 6 2 4 6 2 2 4 4 2 4 ...
rm(file_url, csv_filepath)
```

To improve the domain knowledge, you can find below an image illustrating the different parts of a mushroom:
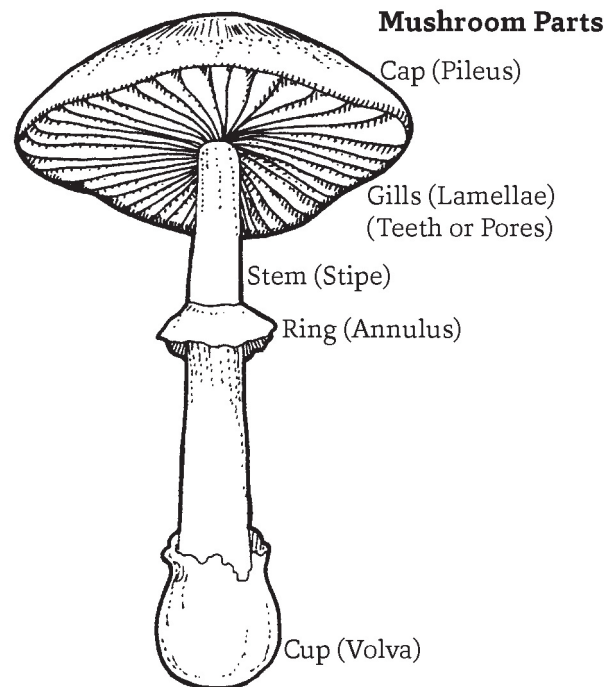
All the attributes are factors and they represent the following abbreviations:

1. cap-shape: bell=b,conical=c,convex=x,flat=f, knobbed=k,sunken=s
2. cap-surface: fibrous=f,grooves=g,scaly=y,smooth=s
3. cap-color: brown=n,buff=b,cinnamon=c,gray=g,green=r, pink=p,purple=u,red=e,white=w,yellow=y
4. bruises?: bruises=t,no=f
5. odor: almond=a,anise=l,creosote=c,fishy=y,foul=f, musty=m,none=n,pungent=p,spicy=s
6. gill-attachment: attached=a,descending=d,free=f,notched=n
7. gill-spacing: close=c,crowded=w,distant=d
8. gill-size: broad=b,narrow=n
9. gill-color: black=k,brown=n,buff=b,chocolate=h,gray=g, green=r,orange=o,pink=p,purple=u,red=e, white=w,yellow=y
10. stalk-shape: enlarging=e,tapering=t
11. stalk-root: bulbous=b,club=c,cup=u,equal=e, rhizomorphs=z,rooted=r,missing=?
12. stalk-surface-above-ring: fibrous=f,scaly=y,silky=k,smooth=s
13. stalk-surface-below-ring: fibrous=f,scaly=y,silky=k,smooth=s
14. stalk-color-above-ring: brown=n,buff=b,cinnamon=c,gray=g,orange=o, pink=p,red=e,white=w,yellow=y
15. stalk-color-below-ring: brown=n,buff=b,cinnamon=c,gray=g,orange=o, pink=p,red=e,white=w,yellow=y

16. veil-type: partial=p,universal=u
17. veil-color: brown=n,orange=o,white=w,yellow=y
18. ring-number: none=n,one=o,two=t
19. ring-type: cobwebby=c,evanescent=e,flaring=f,large=l, none=n,pendant=p,sheathing=s,zone=z
20. spore-print-color: black=k,brown=n,buff=b,chocolate=h,green=r, orange=o,purple=u,white=w,yellow=y
21. population: abundant=a,clustered=c,numerous=n, scattered=s,several=v,solitary=y
22. habitat: grasses=g,leaves=l,meadows=m,paths=p, urban=u,waste=w,woods=d

The classes used for the outcome are `edible` or `poisonous`. In the Kaggle version of the data set there is no `unknown` classification value and missing values have been removed:

```
unique(mushrooms$class)
```

```
## [1] p e
## Levels: e p
```

```
sum(is.na(mushrooms))
```

```
## [1] 0
```

Analysis

The main challenges raised by this data set are:

a. all the features are categorical => the core of the solution will be in the data wrangling part
b. the amount of observation is small (overall around 8k, but we will split it into a training and test sets) => we will need to implement some techniques to avoid overfitting

Data Wrangling

If we try to approach this classification problem as the `tissue_gene_expression` data set distributed by the `dslabs` package, we will quickly realise that the most of utilities used in the PH125.x courses won't be available.

```
# format as list of a features matrix and outcome vector
formatted_mushrooms <- list(
  x = mushrooms %>% select(-class) %>% as.matrix,
  y = mushrooms$class
)

# Partition train and test set
set.seed(22, sample.kind="Rounding")
test_index <- createDataPartition(y=formatted_mushrooms$y, times=1, p=.30, list=FALSE)

train_set <- list(
  x = formatted_mushrooms$x[-test_index,],
  y = formatted_mushrooms$y[-test_index]
)

hclust(train_set$x)
# Error in if (is.na(n) || n > 65536L) stop("size cannot be NA nor exceed 65536") :
# missing value where TRUE/FALSE needed

pca <- prcomp(train_set$x)
# Error in colMeans(x, na.rm = TRUE) : 'x' must be numeric
```

The problem is the most of libraries cannot work with categorical features only. A solution to this problem can be *encoding the categorical features*. Even if the most of resources I googled were explaining how to perform this technique with python (All about Categorical Variable Encoding, How to Encode Categorical Data), there is an excellent R blog post providing a wide overview on the topic.

The most important concept to recall before continuing are:

- Categorical variables can be considered:
    - *Nominal* (e.g. pen/pencil/eraser or cow/dog/cat)
    - *Ordinal* (e.g. excellent/good/bad or fantastic/ok/don't like).

- The three main routes to encode factors string data type are:
    - *Classic Encoders*: e.g. ordinal, OneHot, Binary, Frequency, Hashing ...
    - *Contrast Encoders*: encode data by looking at different levels of features
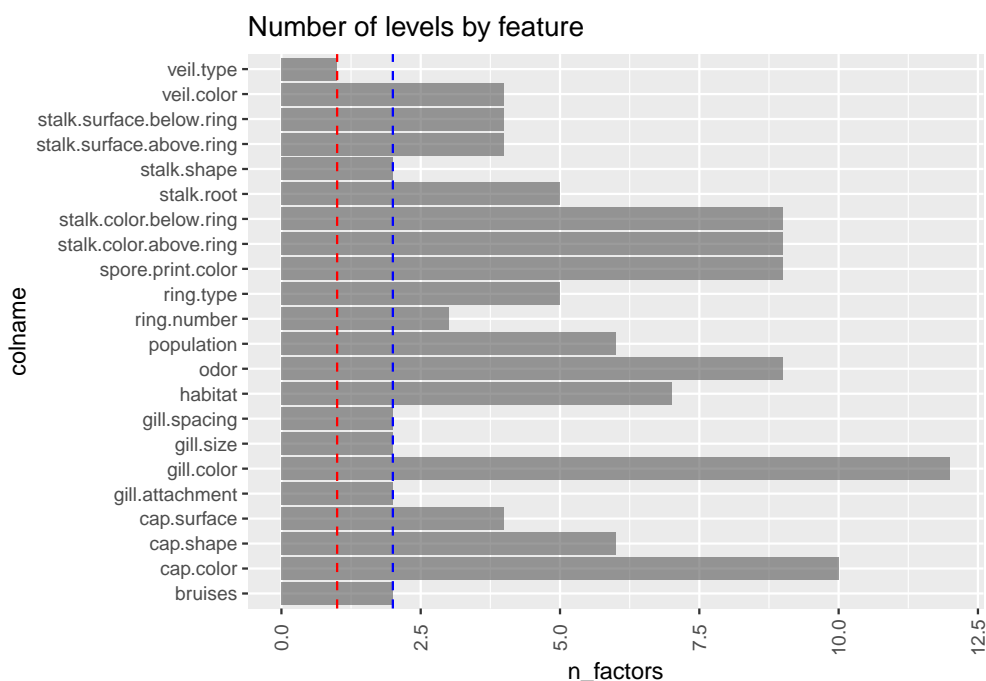   - *Bayesian Encoders*: use the target as the foundation of the encoding

Table 1: Attributes with less than 4 factors

colname	n_factors
ring.number	3
bruises	2
gill.attachment	2
gill.spacing	2
gill.size	2
stalk.shape	2
veil.type	1

We can make these insights actionable with the attributes description provided in the Introduction:

- `veil.type` attribute can be ignored since it has only one level
- `bruises`, `gill.size` attributes can be considered as a logical/binary vectors
- `ring.number` attribute is an ordinal variable

- `population` attribute may be considered as an ordinal variable as well, but it is ambigous (at least for my domain understanding level)
- the other 18 attributes may be considered as nominal variables

Exploration without any normalization

Before formatting the data, we can only explore the data manually. In fact, we cannot use any library we got familiar with.

Some insights were already gained:

a. Odor effect

```r
mushrooms %>%
  group_by(odor) %>%
  summarize(n=n(), edible=mean(class=="e")) %>%
  mutate(is_edible=ifelse(
        edible==1, "yes", ifelse(edible==0, "no", "maybe"))
  ) %>%
  ggplot(aes(x=odor), group=1) +
  geom_bar(aes(y=n, fill=is_edible), stat="identity") +
  geom_text(aes(y=n, label=paste(100*round(edible, digits=3), "%", sep="")), position=position_dodge(wid
  scale_fill_manual(values=c("#f7ef99", "#f78e69","#afc97e")) +
  scale_x_discrete(labels=c(
    "a"="almond","l"="anise","c"="creosote","y"="fishy","f"="foul",
    "m"="musty","n"="none","p"="pungent","s"="spicy"
  )) +
  ggtitle("Mushrooms  odor effect") +
  theme(axis.text.x = element_text(angle=45))
```


Mushrooms odor effect
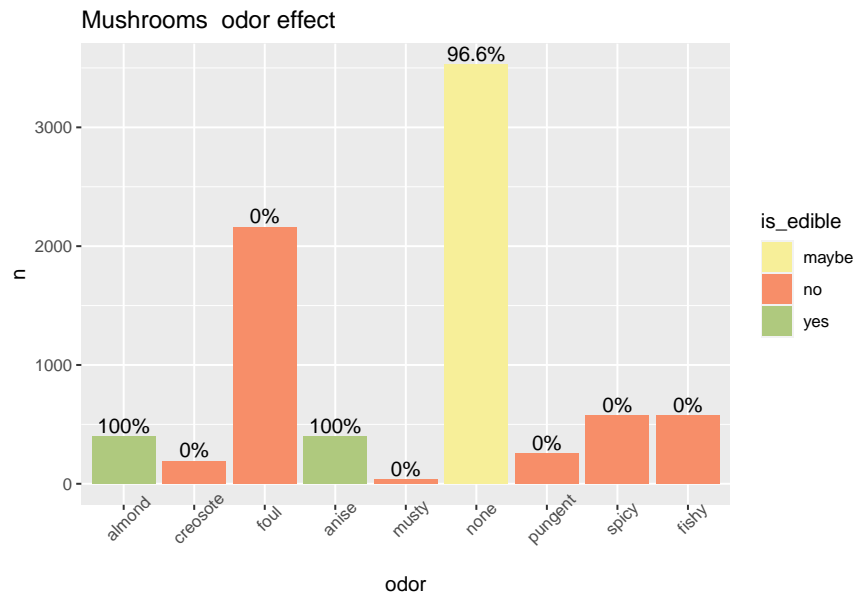
- When odor is almond or anise, the mushroom is alway edible.
- When the mushroom has no odor, it has a 96.6% probability to be edible.
- In any other case, the mushroom s not edible.

It would be worth then to explore investigate the effect of the other features when odor is none.

b. Cap.color effect

```r
mushrooms %>%
  filter(odor =="n") %>%
  group_by(cap.color) %>%
  summarize(n=n(), edible=mean(class=="e")) %>%
  mutate(is_edible=ifelse(
       edible==1, "yes", ifelse(edible==0, "no", "maybe"))
  ) %>%
  ggplot(aes(x=cap.color), group=1) +
  geom_bar(aes(y=n, fill=is_edible), stat="identity") +
  geom_text(aes(y=n, label=paste(100*round(edible, digits=3), "%", sep="")), position=position_dodge(wid
  scale_fill_manual(values=c("#f7ef99", "#f78e69","#afc97e")) +
  scale_x_discrete(labels=c(
    "n"="brown","b"="buff","c"="cinnamon","g"="gray","r"="green",
    "p"="pink","u"="purple","e"="red","w"="white","y"="yellow"
  )) +
  ggtitle("Mushrooms cap color effect with none odor") +
  theme(axis.text.x = element_text(angle=45))
```
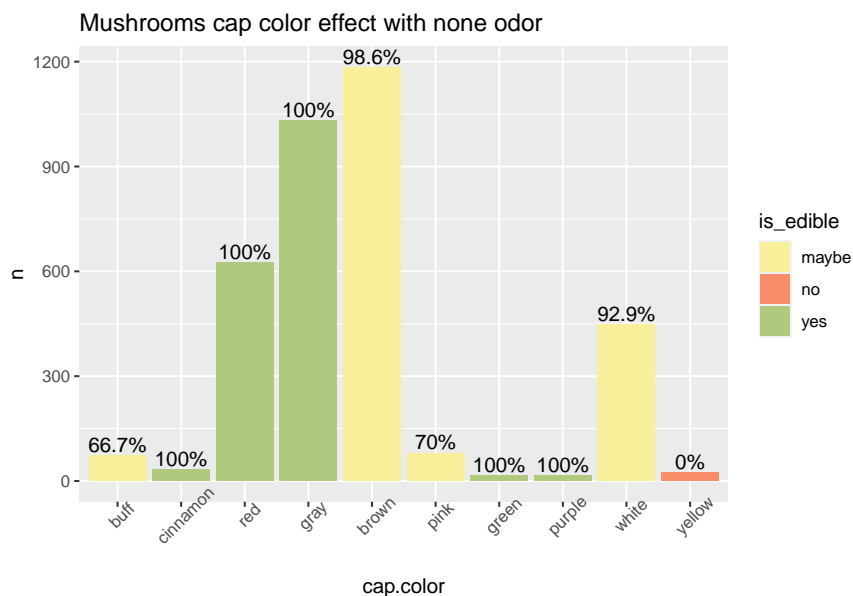

The additional insights analysing the mushrooms with odor `none` by cap color are:

- When cap is cinnamon, red, gray, green or purple the mushrooms are edible
- When cap is yellow, the mushroom is never edible

c. Gill.color effect

```r
mushrooms %>%
  filter(odor =="n") %>%
  filter(cap.color %in% c("n", "b", "p", "w")) %>%
  group_by(gill.color) %>%
  summarize(n=n(), edible=mean(class=="e")) %>%
  mutate(is_edible=ifelse(
       edible==1, "yes", ifelse(edible==0, "no", "maybe"))
  ) %>%
  ggplot(aes(x=gill.color), group=1) +
  geom_bar(aes(y=n, fill=is_edible), stat="identity") +
```

```r
  geom_text(aes(y=n, label=paste(100*round(edible, digits=3), "%", sep="")), position=position_dodge(wi
  scale_fill_manual(values=c("#f7ef99", "#f78e69","#afc97e")) +
  scale_x_discrete(labels=c(
    "k"="black","n"="brown","b"="buff","h"="chocolate","g"="gray",
    "r"="green","o"="orange","p"="pink","u"="purple","e"="red",
    "w"="white","y"="yellow"
  )) +
  ggtitle("Gill color effect on a subset of mushrooms") +
  theme(axis.text.x = element_text(angle=45))
```
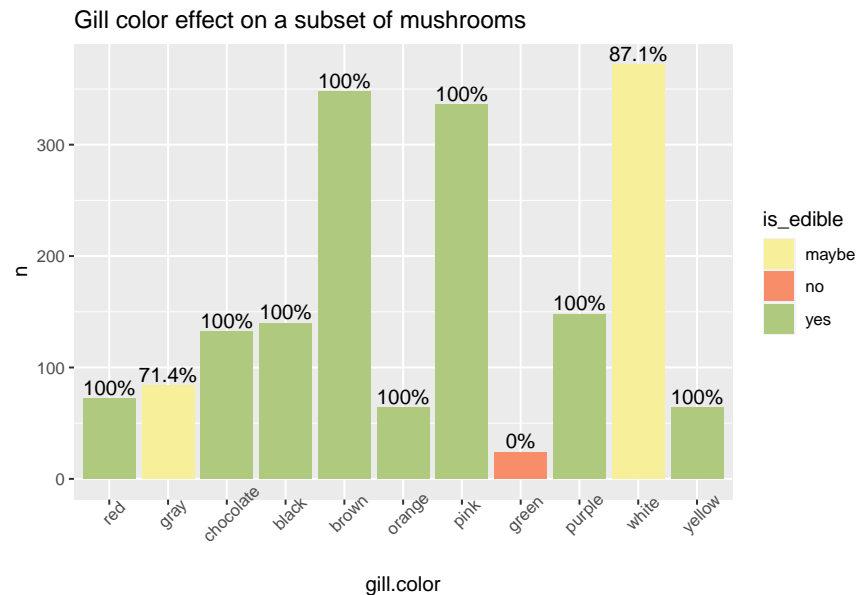

The insights gained on mushroom with none odor and a color brown, buff, pink or white are:

- When gill.color is red, chocolate, black, brown, orange, pink, purple or yellow the mushroom is always edible
- When the gill.color is green the mushroom is not edible

We could have continue this breakdown to build a model to validate our predictions against the true values. We can use however a set of techniques to manipulate categorical attributes as described above.

Using Classic Encoders

A different (and still easy to understand) approach would be to encode these categorical features using **OneHot** encoding.

A one hot encoding is a representation of categorical variables as binary vectors. This can be perfomed using some libraries as vtreat as follows:

```r
if(!require(vtreat)) install.packages("vtreat", repos = "http://cran.r-project.org")
tz <- vtreat::designTreatmentsZ(x, colnames(x))
new_x <- vtreat::prepare(tz, x)
ncol(new_x)
colnames(new_x)
detach("package:vtreat", unload = TRUE)
```

The output of this normalization is very effective for creating efficient models but it may lead to some lack of visibility when interpreting the results.

In our analysis, we are going to normalize the data manually using different techniques before splitting train and test set. In this way we can keep a strict control on the formatted data structure. In order to make this manipulation easy to reproduce (e.g. in case we want to test our model against additional observations), it will be wrapped in a function called `formatMushrooms`:

```r
formatMushrooms <- function(df) {
  # private utility function
  mapValues <- function(v, m) {
    sapply(v, function(x) do.call("switch", prepend(m, x)))
  }

  # Select nominal categorical features with more than 2 categories
  new_df <- df %>%
    select(
      cap.color, cap.shape, cap.surface, gill.attachment, gill.color, habitat, odor,
      population, ring.type, spore.print.color, stalk.color.above.ring,
      stalk.color.below.ring, stalk.root, stalk.surface.above.ring,
      stalk.surface.below.ring, veil.color
    )

  # Nominal variables
  colorsMapping <- list(
    "b"=".buff","c"=".cinnamon","e"=".red","g"=".gray","h"=".chocolate","k"=".black",
    "n"=".brown","o"=".orange","p"=".pink","r"=".green","u"=".purple","w"=".white",
    "y"=".yellow"
  )
  surfacesMapping <- list("f"=".fibrous","g"=".grooves","y"=".scaly","s"=".smooth")

  capShapeMapping <- list("b"=".bell","c"=".conical","x"=".convex","f"=".flat", "k"=".knobbed","s"=".su
  gillAttachmentMapping <- list("a"=".attached","d"=".descending","f"=".free","n"=".notched")
  habitatMapping <- list("g"=".grasses","l"=".leaves","m"=".meadows","p"=".paths", "u"=".urban","w"=".wa
  odorMapping <- list("a"=".almond","l"=".anise","c"=".creosote","y"=".fishy","f"=".foul", "m"=".musty"
  populationMapping <- list("a"=".abundant","c"=".clustered","n"=".numerous", "s"=".scattered","v"=".sev
  ringTypeMapping <- list("c"=".cobwebby","e"=".evanescent","f"=".flaring","l"=".large", "n"=".none","p
  stalkRootMapping <- list("b"=".bulbous","c"=".club","u"=".cup","e"=".equal", "z"=".rhizomorphs","r"="

  new_df$cap.color <- mapValues(df$cap.color, colorsMapping) %>% factor
  new_df$gill.color <- mapValues(df$gill.color, colorsMapping) %>% factor
  new_df$spore.print.color <- mapValues(df$spore.print.color, colorsMapping) %>% factor
  new_df$stalk.color.above.ring <- mapValues(df$stalk.color.above.ring, colorsMapping) %>% factor
  new_df$stalk.color.below.ring <- mapValues(df$stalk.color.below.ring, colorsMapping) %>% factor
  new_df$veil.color <- mapValues(df$veil.color, colorsMapping) %>% factor

  new_df$cap.surface <- mapValues(df$cap.surface, surfacesMapping) %>% factor
  new_df$stalk.surface.above.ring <- mapValues(df$stalk.surface.above.ring, surfacesMapping) %>% factor
  new_df$stalk.surface.below.ring <- mapValues(df$stalk.surface.below.ring, surfacesMapping) %>% factor

  new_df$cap.shape <- mapValues(df$cap.shape, capShapeMapping) %>% factor
  new_df$gill.attachment <- mapValues(df$gill.attachment, gillAttachmentMapping) %>% factor
  new_df$habitat <- mapValues(df$habitat, habitatMapping) %>% factor
  new_df$odor <- mapValues(df$odor, odorMapping) %>% factor
  new_df$population <- mapValues(df$population, populationMapping) %>% factor
  new_df$ring.type <- mapValues(df$ring.type, ringTypeMapping) %>% factor
  new_df$stalk.root <- mapValues(df$stalk.root, stalkRootMapping) %>% factor
```

```r
  # Cleanup
  rm(
    colorsMapping, surfacesMapping, capShapeMapping, gillAttachmentMapping,
    habitatMapping, odorMapping, populationMapping, ringTypeMapping, stalkRootMapping
  )
  # One Hot encoding: convert nominal variables with more than 2 categories
  # https://www.rdocumentation.org/packages/stats/versions/3.6.2/topics/model.matrix
  new_df <- model.matrix(
    ~.-1,
    data = new_df
  )

  # ordinal variables
  # gill.spacing: "c"="close","w"="crowded","d"="distant"
  gillSpacingMapping <- list("c"=0,"w"=1,"d"=2)
  # ring.number: "n"="none","o"="one","t"="two"
  ringNumberMapping <-list("n"=0,"o"=1,"t"=2)

  new_df <- new_df %>%
    as_tibble() %>%
    add_column(
      grill.spacing = mapValues(df$gill.spacing, gillSpacingMapping),
      ring.number = mapValues(df$ring.number, ringNumberMapping),
      bruises = ifelse(df$bruises=="t",1,0),
      gill.size.narrow = ifelse(df$gill.size=="n",1,0),
      stalk.shape.enrlarging = ifelse(df$stalk.shape=="e",1,0)
      # veil.type is ignored in this analysis since it is constant in the data set
      # veil.type.universal = ifelse(df$veil.type=="u",1,0)
    ) %>% as.matrix

  rm(gillSpacingMapping, ringNumberMapping)

  list(
    x = new_df,
    y = df$class == "e" # edible
  )
}
```

The formatted data are:

```r
formatted_mushrooms <- formatMushrooms(mushrooms)
rm(mushrooms)
```